

# Tidy data & dplyr

## Lecture 06

Dr. Colin Rundel



# Tidy data

country	year	cases	population
Afghanistan	1999	76745	19987071
Afghanistan	2000	8666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	216766	128042583

variables

country	year	cases	population
Afghanistan	1999	76745	19987071
Afghanistan	2000	8666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	216766	128042583

observations

country	year	cases	population
Afghanistan	1999	76745	19987071
Afghanistan	2000	8666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	216766	128042583

values



# Tidy vs Untidy

Happy families are all alike; every unhappy family is unhappy in its own way

— Leo Tolstoy, Anna Karenina

# A tibble: 317 × 7

	artist	track	date.entered	wk1	wk2	wk3	wk4
	<chr>	<chr>	<date>	<dbl>	<dbl>	<dbl>	<dbl>
1	2 Pac	Baby Don't Cry (Kee...	2000-02-26	87	82	72	77
2	2Ge+her	The Hardest Part Of...	2000-09-02	91	87	92	NA
3	3 Doors Down	Kryptonite	2000-04-08	81	70	68	67
4	3 Doors Down	Loser	2000-10-21	76	76	72	69
5	504 Boyz	Wobble Wobble	2000-04-15	57	34	25	17
6	98^0	Give Me Just One Ni...	2000-08-19	51	39	34	26
7	A*Teens	Dancing Queen	2000-07-08	97	97	96	95
8	Aaliyah	I Don't Wanna	2000-01-29	84	62	51	41
9	Aaliyah	Try Again	2000-03-18	59	53	38	28
10	Adams Volande	Open My Heart	2000-08-26	76	76	74	60

Is the above data set tidy?

# More tidy vs untidy

Is the following data tidy?

List of 3

\$ :List of 8

```
..$ name      : chr "Luke Skywalker"  
..$ height    : chr "172"  
..$ mass      : chr "77"  
..$ hair_color: chr "blond"  
..$ skin_color: chr "fair"  
..$ eye_color : chr "blue"  
..$ birth_year: chr "19BBY"  
..$ gender    : chr "male"
```

\$ :List of 8

```
..$ name      : chr "C-3PO"  
..$ height    : chr "167"
```

List of 3

\$ :List of 8

```
..$ name      : chr "Darth Vader"  
..$ height    : chr "202"  
..$ mass      : chr "136"  
..$ hair_color: chr "none"  
..$ skin_color: chr "white"  
..$ eye_color : chr "yellow"  
..$ birth_year: chr "41.9BBY"  
..$ gender    : chr "male"
```

\$ :List of 8

```
..$ name      : chr "Leia Organa"  
..$ height    : chr "150"
```





# Modern data frames

The tidyverse includes the tibble package that extends data frames to be a bit more modern. The core features of tibbles is to have a nicer printing method as well as being “surly” and “lazy”.

```
1 library(tibble)
```

```
1 iris
```

```
1 (tbl_iris = as_tibble(iris))
```

	Sepal.Length	Sepal.Width	Petal.Length
1	5.1	3.5	1.4
2	4.9	3.0	1.4
3	4.7	3.2	1.3
4	4.6	3.1	1.5
5	5.0	3.6	1.4
6	5.4	3.9	1.7
7	4.6	3.4	1.4
8	5.0	3.4	1.5
9	4.4	2.9	1.4
10	4.9	3.1	1.5
11	5.4	3.7	1.5
12	4.8	3.4	1.6
13	4.8	3.0	1.4
14	4.3	3.0	1.1
15	5.8	4.0	1.2
16	5.7	4.4	1.5

```
# A tibble: 150 × 5
  Sepal.Length Sepal.Width Petal.Length
      <dbl>         <dbl>         <dbl>
1         5.1         3.5         1.4
2         4.9         3         1.4
3         4.7         3.2         1.3
4         4.6         3.1         1.5
5          5          3.6         1.4
6         5.4         3.9         1.7
7         4.6         3.4         1.4
8          5          3.4         1.5
9         4.4         2.9         1.4
10        4.9         3.1         1.5
# i 140 more rows
# i 2 more variables: Petal.Width <dbl>,
#   Species <fct>
```

# Tibbles are lazy

By default, subsetting tibbles always results in another tibble (`$` or `[[` can still be used to subset for a specific column). i.e. tibble subsets are always preserving and therefore type consistent.

```
1 tbl_iris[1,]
```

```
# A tibble: 1 × 5
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
	<dbl>	<dbl>	<dbl>	<dbl>	<fct>
1	5.1	3.5	1.4	0.2	setosa

```
1 tbl_iris[,1]
```

```
# A tibble: 150 × 1
```

```
  Sepal.Length
```

```
    <dbl>
```

```
1      5.1
2      4.9
3      4.7
4      4.6
5       5
6      5.4
7      4.6
8       5
9      4.4
10     4.9
```

```
# i 140 more rows
```

```
1 head(tbl_iris[[1]])
```

```
[1] 5.1 4.9 4.7 4.6 5.0 5.4
```

```
1 head(tbl_iris$Species)
```

```
[1] setosa setosa setosa setosa setosa setosa
```

```
Levels: setosa versicolor virginica
```

# No partial matching

Tibbles do not use partial matching when the `$` operator is used.

```
1 head( iris$Species )
```

```
[1] setosa setosa setosa setosa setosa  
setosa
```

```
Levels: setosa versicolor virginica
```

```
1 head( iris$Sp )
```

```
[1] setosa setosa setosa setosa setosa  
setosa
```

```
Levels: setosa versicolor virginica
```

```
1 head( tbl_iris$Species )
```

```
[1] setosa setosa setosa setosa setosa  
setosa
```

```
Levels: setosa versicolor virginica
```

```
1 head( tbl_iris$Sp )
```

```
NULL
```

# Tibbles and length coercion

Only vectors with length 1 will undergo length coercion / recycling - everything else will throw an error.

```
1 data.frame(x = 1:4, y = 1)
```

	x	y
1	1	1
2	2	1
3	3	1
4	4	1

```
1 tibble(x = 1:4, y = 1)
```

```
# A tibble: 4 × 2
      x     y
  <int> <dbl>
1     1     1
2     2     1
3     3     1
4     4     1
```

```
1 data.frame(x = 1:4, y = 1:2)
```

	x	y
1	1	1
2	2	2
3	3	1
4	4	2

```
1 tibble(x = 1:4, y = 1:2)
```

```
Error in `tibble()`:
! Tibble columns must have compatible
sizes.
• Size 4: Existing data.
• Size 2: Column `y`.
i Only values of size one are recycled.
```

# Tibbles and S3

```
1 t = tibble(  
2   x = 1:3,  
3   y = c("A", "B", "C")  
4 )  
5  
6 class(t)
```

```
[1] "tbl_df"      "tbl"        "data.frame"
```

```
1 d = data.frame(  
2   x = 1:3,  
3   y = c("A", "B", "C")  
4 )  
5  
6 class(d)
```

```
[1] "data.frame"
```

```
1 methods(class="tbl_df")
```

```
[1] [           [[           [[<-          [<-          $  
[6] $<-          as.data.frame coerce       initialize    names<-  
[11] Ops           row.names<-   show         slotsFromS3   str  
[16] tbl_sum
```

see '?methods' for accessing help and source code

```
1 methods(class="tbl")
```

```
[1] [[<-          [<-          $<-          coerce       format  
[6] glimpse       initialize  Ops           print         show  
[11] slotsFromS3  tbl_sum
```

see '?methods' for accessing help and source code

# Supporting tibbles?

```
1 d = tibble(  
2   x = rnorm(100),  
3   y = 3 + x + rnorm(100, sd = 0.1)  
4 )
```

```
1 lm(y~x, data = d)
```

Call:

```
lm(formula = y ~ x, data = d)
```

Coefficients:

(Intercept)	x
3.0058	0.9831

## Why did this work?



magrittr

Sta 323 - Spring 2024



# What is a pipe

In software engineering, a pipeline consists of a chain of processing elements (processes, threads, coroutines, functions, etc.), arranged so that the output of each element is the input of the next;

[Wikipedia - Pipeline \(software\)](#)

Magrittr's pipe is a new infix operator that allows us to link two functions together in a way that is readable from left to right.

The two code examples below are equivalent,

```
1 f(g(x=1, y=2), n=2)
```

```
1 g(x=1, y=2) %>% f(n=2)
```

# Readability

Consider the following sequence of actions that describe the process of getting to campus in the morning:

I need to find my key, then unlock my car, then start my car, then drive to school, then park.

Expressed as a set of nested functions in R pseudocode this would look like:

```
1 park(drive(start_car(find("keys")), to="campus"))
```

Writing it out using pipes give it a more natural (and easier to read) structure:

```
1 find("keys") %>%  
2   start_car() %>%  
3   drive(to="campus") %>%  
4   park()
```

# Approaches

All of the following are fine, it comes down to personal preference:

Nested:

```
1 h( g( f(x), y=1 ), z=1 )
```

Piped:

```
1 f(x) %>%  
2   g(y=1) %>%  
3   h(z=1)
```

Intermediate:

```
1 res = f(x)  
2 res = g(res, y=1)  
3 res = h(res, z=1)
```

# What about other arguments?

Sometimes we want to send our results to an function argument other than first one or we want to use the previous result for multiple arguments. In these cases we can refer to the previous result using `..`

```
1 data.frame(a = 1:3, b = 3:1) %>% lm(a~b, data=.)
```

Call:

```
lm(formula = a ~ b, data = .)
```

Coefficients:

(Intercept)	b
4	-1

```
1 data.frame(a = 1:3, b = 3:1) %>% .[[1]]
```

```
[1] 1 2 3
```

```
1 data.frame(a = 1:3, b = 3:1) %>% .[[length(.)]]
```

```
[1] 3 2 1
```

# The base R pipe

As of R v4.1.0 a native pipe operator was added to the base language in R, it is implemented as `|>`.

```
1 1:10 |> cumsum()
```

```
[1] 1 3 6 10 15 21 28 36 45 55
```

```
1 1:10 |> cumsum() |> mean()
```


```
[1] 22
```

The current version of RStudio on the departmental servers is v4.3.2 so you are welcome to try it out.

# Base R pipe considerations:

- Depending on R version  $\geq 4.1$  is a harder dependency than depending on the magrittr package
- `|>` will likely have less overhead than `%>%` but the difference is unlikely to matter in practice
- `|>` supports an equivalent to `.` using `_` as of R v4.2 (but only for named arguments)

Generally we will prefer the base pipe in this class, but using either is fine.



# A Grammar of Data Manipulation

dplyr is based on the concepts of functions as verbs that manipulate data frames.

Core single data frame functions / verbs:

- `filter()` / `slice()`: pick rows based on criteria
- `select()` / `rename()`: select columns by name
- `pull()`: grab a column as a vector
- `arrange()`: reorder rows
- `mutate()` / `transmute()`: create or modify columns
- `distinct()`: filter for unique rows
- `summarise()` / `count()`: reduce variables to values
- `group_by()` / `ungroup()`: modify other verbs to act on subsets
- `relocate()`: change column order
- ... (many more)



# dplyr mental model / rules

1. First argument is *always* a data frame
2. Subsequent arguments say what to do with the data frame
3. *Always* return a data frame
4. Don't modify in place
5. Magic via lazy evaluation and S3

# Example Data

We will demonstrate dplyr's functionality using the nycflights13 data.

```
1 library(dplyr)
2 library(nycflights13)
```

```
1 flights
```

```
# A tibble: 336,776 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	-2	753

```
# i 336,766 more rows
```

# filter() - March flights

```
1 flights |> filter(month == 3)
```

```
# A tibble: 28,834 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	3	1	4	2159	125	318
2	2013	3	1	50	2358	52	526
3	2013	3	1	117	2245	152	223
4	2013	3	1	454	500	-6	633
5	2013	3	1	505	515	-10	746
6	2013	3	1	521	530	-9	813
7	2013	3	1	537	540	-3	856
8	2013	3	1	541	545	-4	1014
9	2013	3	1	549	600	-11	639
10	2013	3	1	550	600	10	717

# filter() - Flights in the first 7 days of March

```
1 flights |> filter(month == 3, day <= 7)
```

```
# A tibble: 6,530 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	3	1	4	2159	125	318
2	2013	3	1	50	2358	52	526
3	2013	3	1	117	2245	152	223
4	2013	3	1	454	500	-6	633
5	2013	3	1	505	515	-10	746
6	2013	3	1	521	530	-9	813
7	2013	3	1	537	540	-3	856
8	2013	3	1	541	545	-4	1014
9	2013	3	1	549	600	-11	639
10	2013	3	1	550	600	10	717

# filter() - Flights to LAX or JFK in March

```
1 flights |> filter(dest == "LAX" | dest == "JFK", month==3)
```

```
# A tibble: 1,178 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	3	1	607	610	-3	832
2	2013	3	1	629	632	-3	844
3	2013	3	1	657	700	-3	953
4	2013	3	1	714	715	-1	939
5	2013	3	1	716	710	6	958
6	2013	3	1	727	730	-3	1007
7	2013	3	1	836	840	-4	1111
8	2013	3	1	857	900	-3	1202
9	2013	3	1	903	900	3	1157
10	2013	3	1	904	901	3	1158

# slice() - First 10 flights

```
1 flights |> slice(1:10)
```

```
# A tibble: 10 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	2	753

# slice() - Last 5 flights

```
1 flights |> slice((n()-4):n())
```

```
# A tibble: 5 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	9	30	NA	1455	NA	NA
2	2013	9	30	NA	2200	NA	NA
3	2013	9	30	NA	1210	NA	NA
4	2013	9	30	NA	1159	NA	NA
5	2013	9	30	NA	840	NA	NA

```
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,  
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
#   time_hour <dtm>
```

# slice\_tail() - Last 5 flights

```
1 flights |> slice_tail(n = 5)
```

```
# A tibble: 5 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	9	30	NA	1455	NA	NA
2	2013	9	30	NA	2200	NA	NA
3	2013	9	30	NA	1210	NA	NA
4	2013	9	30	NA	1159	NA	NA
5	2013	9	30	NA	840	NA	NA

```
# i 12 more variables: sched_arr_time <int>, arr_delay <dbl>,  
#   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
#   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
#   time_hour <dtm>
```



# select() - Individual Columns

```
1 flights |> select(year, month, day)
```

```
# A tibble: 336,776 × 3
```

	year	month	day
	<int>	<int>	<int>
1	2013	1	1
2	2013	1	1
3	2013	1	1
4	2013	1	1
5	2013	1	1
6	2013	1	1
7	2013	1	1
8	2013	1	1
9	2013	1	1
10	2013	1	1

# select() - Exclude Columns

```
1 flights |> select(-year, -month, -day)
```

# A tibble: 336,776 × 16

	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time	arr_delay
	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	517	515	2	830	819	11
2	533	529	4	850	830	20
3	542	540	2	923	850	33
4	544	545	-1	1004	1022	-18
5	554	600	-6	812	837	-25
6	554	558	-4	740	728	12
7	555	600	-5	913	854	19
8	557	600	-3	709	723	-14
9	557	600	-3	838	846	-8
10	558	600	2	753	745	8

# select() - Ranges

```
1 flights |> select(year:day)
```

```
# A tibble: 336,776 × 3
```

	year	month	day
	<int>	<int>	<int>
1	2013	1	1
2	2013	1	1
3	2013	1	1
4	2013	1	1
5	2013	1	1
6	2013	1	1
7	2013	1	1
8	2013	1	1
9	2013	1	1
10	2013	1	1

# select() - Exclusion Ranges

```
1 flights |> select(-(year:day))
```

```
# A tibble: 336,776 × 16
```

	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time	arr_delay
	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	517	515	2	830	819	11
2	533	529	4	850	830	20
3	542	540	2	923	850	33
4	544	545	-1	1004	1022	-18
5	554	600	-6	812	837	-25
6	554	558	-4	740	728	12
7	555	600	-5	913	854	19
8	557	600	-3	709	723	-14
9	557	600	-3	838	846	-8
10	558	600	2	753	745	8

# select() - Matching contains()

```
1 flights |> select(contains("dep"), contains("arr"))
```

```
# A tibble: 336,776 × 7
```

	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time	arr_delay
	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	517	515	2	830	819	11
2	533	529	4	850	830	20
3	542	540	2	923	850	33
4	544	545	-1	1004	1022	-18
5	554	600	-6	812	837	-25
6	554	558	-4	740	728	12
7	555	600	-5	913	854	19
8	557	600	-3	709	723	-14
9	557	600	-3	838	846	-8
10	558	600	2	753	745	8

# select() - Matching starts\_with()

```
1 flights |> select(starts_with("dep"), starts_with("arr"))
```

```
# A tibble: 336,776 × 4
```

	dep_time	dep_delay	arr_time	arr_delay
	<int>	<dbl>	<int>	<dbl>
1	517	2	830	11
2	533	4	850	20
3	542	2	923	33
4	544	-1	1004	-18
5	554	-6	812	-25
6	554	-4	740	12
7	555	-5	913	19
8	557	-3	709	-14
9	557	-3	838	-8
10	558	0	753	0

Other helpers provide by [tidyselect](#):

# select() + where() - Get numeric columns

```
1 flights |> select(where(is.numeric))
```

```
# A tibble: 336,776 × 14
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	-2	753

```
# i 336,766 more rows
```

# select() + where() - Get non-numeric columns

```
1 flights |> select(where(function(x) !is.numeric(x)))
```

```
# A tibble: 336,776 × 5
```

	carrier	tailnum	origin	dest	time_hour
	<chr>	<chr>	<chr>	<chr>	<dtm>
1	UA	N14228	EWR	IAH	2013-01-01 05:00:00
2	UA	N24211	LGA	IAH	2013-01-01 05:00:00
3	AA	N619AA	JFK	MIA	2013-01-01 05:00:00
4	B6	N804JB	JFK	BQN	2013-01-01 05:00:00
5	DL	N668DN	LGA	ATL	2013-01-01 06:00:00
6	UA	N39463	EWR	ORD	2013-01-01 05:00:00
7	B6	N516JB	EWR	FLL	2013-01-01 06:00:00
8	EV	N829AS	LGA	IAD	2013-01-01 06:00:00
9	B6	N593JB	JFK	MCO	2013-01-01 06:00:00
10	AA	N3ALAA	LGA	ORD	2013-01-01 06:00:00

```
# i 336,766 more rows
```



# relocate - to the front

```
1 flights |> relocate(carrier, origin, dest)
```

# A tibble: 336,776 × 19

	carrier	origin	dest	year	month	day	dep_time	sched_dep_time	dep_delay
	<chr>	<chr>	<chr>	<int>	<int>	<int>	<int>	<int>	<dbl>
1	UA	EWR	IAH	2013	1	1	517	515	2
2	UA	LGA	IAH	2013	1	1	533	529	4
3	AA	JFK	MIA	2013	1	1	542	540	2
4	B6	JFK	BQN	2013	1	1	544	545	-1
5	DL	LGA	ATL	2013	1	1	554	600	-6
6	UA	EWR	ORD	2013	1	1	554	558	-4
7	B6	EWR	FLL	2013	1	1	555	600	-5
8	EV	LGA	IAD	2013	1	1	557	600	-3
9	B6	JFK	MCO	2013	1	1	557	600	-3
10	AA	LGA	ORD	2013	1	1	558	600	-2

# relocate - to the end

```
1 flights |> relocate(year, month, day, .after = last_col())
```

```
# A tibble: 336,776 × 19
```

	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time	arr_delay
	<int>	<int>	<dbl>	<int>	<int>	<dbl>
1	517	515	2	830	819	11
2	533	529	4	850	830	20
3	542	540	2	923	850	33
4	544	545	-1	1004	1022	-18
5	554	600	-6	812	837	-25
6	554	558	-4	740	728	12
7	555	600	-5	913	854	19
8	557	600	-3	709	723	-14
9	557	600	-3	838	846	-8
10	558	600	2	753	745	8

# rename() - Change column names

```
1 flights |> rename(tail_number = tailnum)
```

```
# A tibble: 336,776 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	2	753

# select() vs. rename()

```
1 flights |> select(tail_number = tailnum)
```

```
# A tibble: 336,776 × 1
```

```
  tail_number  
  <chr>
```

```
1 N14228  
2 N24211  
3 N619AA  
4 N804JB  
5 N668DN  
6 N39463  
7 N516JB  
8 N829AS  
9 N593JB  
10 N3ALAA
```

```
# i 336,766 more rows
```

```
1 flights |> rename(tail_number = tailnum)
```

```
# A tibble: 336,776 × 19
```

```
  year month   day dep_time sched_dep_time  
  <int> <int> <int>   <int>         <int>  
1  2013     1     1     517           515  
2  2013     1     1     533           529  
3  2013     1     1     542           540  
4  2013     1     1     544           545  
5  2013     1     1     554           600  
6  2013     1     1     554           558  
7  2013     1     1     555           600  
8  2013     1     1     557           600  
9  2013     1     1     557           600  
10 2013     1     1     558           600
```

```
# i 336,766 more rows
```

```
# i 14 more variables: dep_delay <dbl>,
```

```
#   arr_time <int>, sched_arr_time <int>,
```

```
#   arr_delay <dbl>, carrier <chr>, flight <int>.
```

# pull()

```
1 names(flights)
```

```
[1] "year"          "month"          "day"            "dep_time"
[5] "sched_dep_time" "dep_delay"      "arr_time"       "sched_arr_time"
[9] "arr_delay"      "carrier"        "flight"         "tailnum"
[13] "origin"         "dest"           "air_time"       "distance"
[17] "hour"           "minute"         "time_hour"
```

```
1 flights |> pull("year") |> head()
```

```
[1] 2013 2013 2013 2013 2013 2013
```

```
1 flights |> pull(1) |> head()
```

```
[1] 2013 2013 2013 2013 2013 2013
```

```
1 flights |> pull(-1) |> head()
```

```
[1] "2013-01-01 05:00:00 EST" "2013-01-01 05:00:00 EST"
[3] "2013-01-01 05:00:00 EST" "2013-01-01 05:00:00 EST"
[5] "2013-01-01 06:00:00 EST" "2013-01-01 05:00:00 EST"
```

# arrange() - Sort data

```
1 flights |> filter(month==3,day==2) |> arrange(origin, dest)
```

```
# A tibble: 765 × 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	3	2	1336	1329	7	1426
2	2013	3	2	628	629	-1	837
3	2013	3	2	637	640	-3	903
4	2013	3	2	743	745	-2	945
5	2013	3	2	857	900	-3	1117
6	2013	3	2	1027	1030	-3	1234
7	2013	3	2	1134	1145	-11	1332
8	2013	3	2	1412	1415	-3	1636
9	2013	3	2	1633	1636	-3	1848
10	2013	3	2	1655	1700	5	1857

# arrange() w/ desc() - descending order

```
1 flights |>
2   filter(month==3, day==2) |>
3   arrange(desc(origin), dest) |>
4   select(origin, dest, tailnum)
```

# A tibble: 765 × 3

	origin	dest	tailnum
	<chr>	<chr>	<chr>
1	LGA	ATL	N928AT
2	LGA	ATL	N623DL
3	LGA	ATL	N680DA
4	LGA	ATL	N996AT
5	LGA	ATL	N510MQ
6	LGA	ATL	N663DN
7	LGA	ATL	N942DL
8	LGA	ATL	N511MQ
9	LGA	ATL	N910DE
10	LGA	ATL	N902DE

# distinct() - Find unique rows

```
1 flights |>
2   select(origin, dest) |>
3   distinct() |>
4   arrange(origin,dest)
```

```
# A tibble: 224 × 2
```

	origin	dest
	<chr>	<chr>
1	EWR	ALB
2	EWR	ANC
3	EWR	ATL
4	EWR	AUS
5	EWR	AVL
6	EWR	BDL
7	EWR	BNA
8	EWR	BOS
9	EWR	BQN
10	EWR	DMW



# mutate() - Modify / create columns

```
1 flights |>
2   select(year:day) |>
3   mutate(date = paste(year, month, day, sep="/"))
```

```
# A tibble: 336,776 × 4
```

	year	month	day	date
	<int>	<int>	<int>	<chr>
1	2013	1	1	2013/1/1
2	2013	1	1	2013/1/1
3	2013	1	1	2013/1/1
4	2013	1	1	2013/1/1
5	2013	1	1	2013/1/1
6	2013	1	1	2013/1/1
7	2013	1	1	2013/1/1
8	2013	1	1	2013/1/1
9	2013	1	1	2013/1/1
10	2013	1	1	2013/1/1

# summarise() - Arregate rows

```
1 flights |>
2   summarize(n(), min(dep_delay), max(dep_delay))
```

# A tibble: 1 × 3

	`n()` <int>	`min(dep_delay)` <dbl>	`max(dep_delay)` <dbl>
1	336776	NA	NA

```
1 flights |>
2   summarize(
3     n = n(),
4     min_dep_delay = min(dep_delay, na.rm = TRUE),
5     max_dep_delay = max(dep_delay, na.rm = TRUE)
6   )
```

# A tibble: 1 × 3

	n <int>	min_dep_delay <dbl>	max_dep_delay <dbl>
1	336776	-43	1301

# group\_by()

```
1 flights |> group_by(origin)
```

# A tibble: 336,776 × 19

# Groups: origin [3]

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	3	830

# summarise() with group\_by()

```
1 flights |>
2   group_by(origin) |>
3   summarize(
4     n = n(),
5     min_dep_delay = min(dep_delay, na.rm = TRUE),
6     max_dep_delay = max(dep_delay, na.rm = TRUE)
7   )
```

# A tibble: 3 × 4

	origin	n	min_dep_delay	max_dep_delay
	<chr>	<int>	<dbl>	<dbl>
1	EWR	120835	-25	1126
2	JFK	111279	-43	1301
3	LGA	104662	-33	911

# Groups after summarise

```
1 flights |>
2   group_by(origin) |>
3   summarize(
4     n = n(),
5     min_dep_delay = min(dep_delay, na.rm=TRUE),
6     max_dep_delay = max(dep_delay, na.rm=TRUE),
7     .groups = "drop_last"
8   )
```

# A tibble: 3 × 4

	origin	n	min_dep_delay	max_dep_delay
	<chr>	<int>	<dbl>	<dbl>
1	EWR	120835	-25	1126
2	JFK	111279	-43	1301
3	LGA	104662	-33	911

```
1 flights |>
2   group_by(origin) |>
3   summarize(
4     n = n(),
5     min_dep_delay = min(dep_delay, na.rm=TRUE),
6     max_dep_delay = max(dep_delay, na.rm=TRUE),
7     .groups = "keep"
8   )
```

# A tibble: 3 × 4

# Groups: origin [3]

	origin	n	min_dep_delay	max_dep_delay
	<chr>	<int>	<dbl>	<dbl>
1	EWR	120835	-25	1126
2	JFK	111279	-43	1301
3	LGA	104662	-33	911

# The .by argument

```
1 flights |>
2   summarize(
3     n = n(),
4     min_dep_delay = min(dep_delay, na.rm=TRUE),
5     max_dep_delay = max(dep_delay, na.rm=TRUE),
6     .by = origin
7   )
```

# A tibble: 3 × 4

	origin	n	min_dep_delay	max_dep_delay
	<chr>	<int>	<dbl>	<dbl>
1	EWR	120835	-25	1126
2	LGA	104662	-33	911
3	JFK	111279	-43	1301

The `.by` and `by` arguments are used for per operation grouping while `group_by()` is intended for persistent

# count()

```
1 flights |>
2   summarize(
3     n = n(),
4     .by = c(origin, carrier)
5   )
```

# A tibble: 35 × 3

	origin	carrier	n
	<chr>	<chr>	<int>
1	EWR	UA	46087
2	LGA	UA	8044
3	JFK	AA	13783
4	JFK	B6	42076
5	LGA	DL	23067
6	EWR	B6	6557
7	LGA	EV	8826
8	LGA	AA	15459
9	JFK	UA	4534
10	LGA	B6	6002

```
1 flights |>
2   count(origin, carrier)
```

# A tibble: 35 × 3

	origin	carrier	n
	<chr>	<chr>	<int>
1	EWR	9E	1268
2	EWR	AA	3487
3	EWR	AS	714
4	EWR	B6	6557
5	EWR	DL	4342
6	EWR	EV	43939
7	EWR	MQ	2276
8	EWR	OO	6
9	EWR	UA	46087
10	EWR	US	4405

# mutate() with .by

```
1 flights |>
2   mutate(n = n(), .by = origin) |>
3   select(origin, n)
```

# A tibble: 336,776 × 2

	origin	n
	<chr>	<int>
1	EWR	120835
2	LGA	104662
3	JFK	111279
4	JFK	111279
5	LGA	104662
6	EWR	120835
7	EWR	120835
8	LGA	104662
9	JFK	111279
10	LGA	104662



# Exercises / Examples

1. How many flights to Los Angeles (LAX) did each of the legacy carriers (AA, UA, DL or US) have in May from JFK, and what was their average duration?
2. What was the shortest flight out of each airport in terms of distance? In terms of duration?
3. Which plane (check the tail number) flew out of each New York airport the most?
4. Which date should you fly on if you want to have the lowest possible average departure delay? What about arrival delay?