Tidymodels

Lecture 23

Dr. Colin Rundel
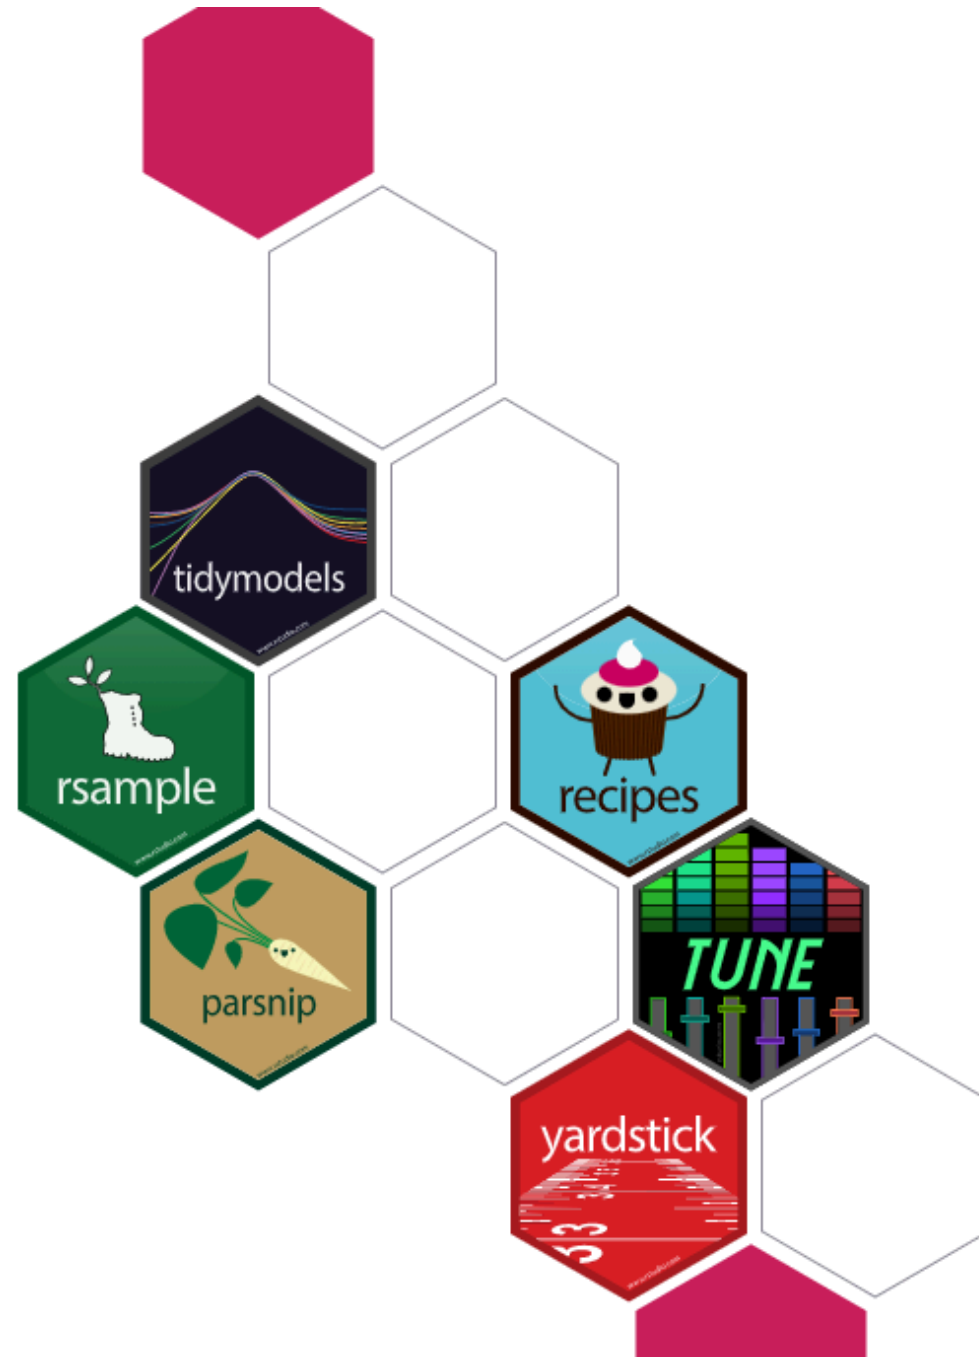
Tidymodels

```
1  library(tidymodels)
```

```
── Attaching packages ──────────────────────────────────── tidymodels 1.2.0 ──
✔ broom        1.0.7      ✔ rsample      1.2.1
✔ dials        1.3.0      ✔ tune         1.2.1
✔ infer        1.0.7      ✔ workflows    1.1.4
✔ modeldata    1.4.0      ✔ workflowsets 1.1.0
✔ parsnip      1.2.1      ✔ yardstick    1.3.1
✔ recipes      1.1.0

── Conflicts ───────────────────────────────────────── tidymodels_conflicts() ──
✖ scales::discard()   masks purrr::discard()
✖ dplyr::filter()     masks stats::filter()
✖ recipes::fixed()    masks stringr::fixed()
✖ dplyr::lag()        masks stats::lag()
✖ rsample::populate() masks Rcpp::populate()
✖ yardstick::spec()   masks readr::spec()
✖ recipes::step()     masks stats::step()
• Learn how to get started at https://www.tidymodels.org/start/
```

Book data

```r
(books = DAAG::allbacks |>
  as_tibble() |>
  select(-area) |>
  mutate(
    cover = forcats::fct_recode(
      cover,
      "hardback" = "hb",
      "paperback" = "pb"
    )
  )
)
```

```
# A tibble: 15 × 3
   volume weight cover
    <dbl>  <dbl> <fct>
 1    885    800 hardback
 2   1016    950 hardback
 3   1125   1050 hardback
 4    239    350 hardback
 5    701    750 hardback
 6    641    600 hardback
 7   1228   1075 hardback
 8    412    250 paperback
 9    953    700 paperback
10    929    650 paperback
```

```
1  ggplot(books, aes(x=volume, y=weight, color = cover)) +
2    geom_point(size=2)
```

Building a tidymodel

```r
1 linear_reg()
```

```
Linear Regression Model Specification (regression)

Computational engine: lm
```

```r
1 linear_reg() |>
2   set_engine("lm")
```

```
Linear Regression Model Specification (regression)

Computational engine: lm
```

Building a tidymodel

```
1  linear_reg() |>
2    set_engine("lm") |>
3    fit(weight ~ volume * cover, data = books)
```

parsnip model object

Call:
stats::lm(formula = weight ~ volume * cover, data =
data)

Coefficients:
       (Intercept)   volume
         161.58654   0.76159
    coverpaperback   volume:coverpaperback
        -120.21407   -0.07573

```
1  lm(weight ~ volume * cover, data = books)
```

Call:
lm(formula = weight ~ volume * cover, data = books)

Coefficients:
       (Intercept)   volume
         161.58654   0.76159
    coverpaperback   volume:coverpaperback
        -120.21407   -0.07573

Tidy model objects

```
1  lm_tm = linear_reg() |>
2    set_engine("lm") |>
3    fit(weight ~ volume * cover, data = books)
```

```
1  lm_b = lm(weight ~ volume * cover, data = books)
```

```
1  summary(lm_tm)
```

```
1  summary(lm_b)
```

	Length	Class	Mode
lvl	0	-none-	NULL
spec	7	linear_reg	list
fit	13	lm	list
preproc	1	-none-	list
elapsed	1	-none-	list
censor_probs	0	-none-	list

```
Call:
lm(formula = weight ~ volume * cover, data = books)

Residuals:
   Min     1Q Median    3Q    Max
-89.67 -32.07 -21.82  17.94 215.91


Coefficients:
                      Estimate Std. Error t value
(Intercept)          161.58654   86.51918   1.868
volume                 0.76159    0.09718   7.837
coverpaperback      -120.21407  115.65899  -1.039
volume:coverpaperback -0.07573    0.12802  -0.592
                     Pr(>|t|)
(Intercept)            0.0887 .
volume               7.94e-06 ***
```

Tidy coefficients

```
1  broom::tidy(lm_tm)
```

```
# A tibble: 4 × 5
  term        estimate std.error statistic p.value
  <chr>          <dbl>     <dbl>     <dbl>   <dbl>
1 (Intercept)  1.62e+2    86.5       1.87  8.87e-2
2 volume       7.62e-1     0.0972    7.84  7.94e-6
3 coverpaper… -1.20e+2   116.       -1.04  3.21e-1
4 volume:cov… -7.57e-2     0.128    -0.592 5.66e-1
```

```
1  broom::tidy(lm_b)
```

```
# A tibble: 4 × 5
  term        estimate std.error statistic p.value
  <chr>          <dbl>     <dbl>     <dbl>   <dbl>
1 (Intercept)  1.62e+2    86.5       1.87  8.87e-2
2 volume       7.62e-1     0.0972    7.84  7.94e-6
3 coverpaper… -1.20e+2   116.       -1.04  3.21e-1
4 volume:cov… -7.57e-2     0.128    -0.592 5.66e-1
```

Tidy statistics

```
1  broom::glance(lm_tm)
```

```
# A tibble: 1 × 12
  r.squared adj.r.squared sigma statistic  p.value    df logLik   AIC   BIC deviance
      <dbl>         <dbl> <dbl>     <dbl>    <dbl> <dbl>  <dbl> <dbl> <dbl>    <dbl>
1     0.930         0.911  80.4      48.5  1.24e-6     3  -84.8  180.  183.   71118.
# ℹ 2 more variables: df.residual <int>, nobs <int>
```

```
1  broom::glance(lm_b)
```

```
# A tibble: 1 × 12
  r.squared adj.r.squared sigma statistic  p.value    df logLik   AIC   BIC deviance
      <dbl>         <dbl> <dbl>     <dbl>    <dbl> <dbl>  <dbl> <dbl> <dbl>    <dbl>
1     0.930         0.911  80.4      48.5  1.24e-6     3  -84.8  180.  183.   71118.
# ℹ 2 more variables: df.residual <int>, nobs <int>
```

Tidy prediction

```
1  broom::augment(lm_tm, new_data = books)
```

```
# A tibble: 15 × 5
   .pred .resid volume weight cover
   <dbl>  <dbl>  <dbl>  <dbl> <fct>
 1  836. -35.6     885    800 hardback
 2  935.  14.6    1016    950 hardback
 3 1018.  31.6    1125   1050 hardback
 4  344.   6.39    239    350 hardback
 5  695.  54.5     701    750 hardback
 6  650. -49.8     641    600 hardback
 7 1097. -21.8    1228   1075 hardback
 8  324. -73.9     412    250 paperback
 9  695.   5.00    953    700 paperback
10  679   28.5     929    650 paperback
```

Putting it together

```
1  lm_tm |>
2    augment(
3      new_data = tidyr::expand_grid(
4        volume = seq(0, 1500, by=5),
5        cover = c("hardback", "paperback") |> as.factor()
6      )
7    ) |>
8    rename(weight = .pred) |>
9    ggplot(aes(x = volume, y = weight, color = cover, group = cover)) +
10     geom_line() +
11     geom_point(data = books)
```

Why do we care?

```
1  show_engines("linear_reg")
```

```
# A tibble: 7 × 2
  engine mode
  <chr>  <chr>
1 lm     regression
2 glm    regression
3 glmnet regression
4 stan   regression
5 spark  regression
6 keras  regression
7 brulee regression
```

```
1  (bayes_tm = linear_reg() |>
2    set_engine(
3      "stan",
4      prior_intercept = rstanarm::student_t(df = 1
5      prior = rstanarm::student_t(df = 1),
6      seed = 1234
7    )
8  )
```

```
Linear Regression Model Specification (regression)

Engine-Specific Arguments:
  prior_intercept = rstanarm::student_t(df = 1)
  prior = rstanarm::student_t(df = 1)
  seed = 1234

Computational engine: stan
```

Fitting with `rstanarm`

```
1  (bayes_tm = bayes_tm |>
2    fit(weight ~ volume * cover, data = books)
3  )
```

Warning: There were 58 divergent transitions after warmup. See
https://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup
to find out why this is a problem and how to eliminate them.

Warning: Examine the pairs() plot to diagnose sampling problems

parsnip model object

stan_glm
 family: gaussian [identity]
 formula: weight ~ volume * cover
 observations: 15
 predictors: 4

                      Median MAD_SD
(Intercept)   95.4   63.9
volume   0.8   0.1
coverpaperback   -0.3   3.6
volume:coverpaperback -0.2   0.1

Auxiliary parameter(s):
      Median MAD_SD
sigma 85.5   18.1

What was actually run?

```
1  linear_reg() |>
2    set_engine(
3      "stan",
4      prior_intercept = rstanarm::student_t(df = 1),
5      prior = rstanarm::student_t(df = 1),
6      seed = 1234
7    ) |>
8    translate()
```

```
Linear Regression Model Specification (regression)

Engine-Specific Arguments:
  prior_intercept = rstanarm::student_t(df = 1)
  prior = rstanarm::student_t(df = 1)
  seed = 1234

Computational engine: stan

Model fit template:
rstanarm::stan_glm(formula = missing_arg(), data = missing_arg(),
    weights = missing_arg(), prior_intercept = rstanarm::student_t(df = 1),
    prior = rstanarm::student_t(df = 1), seed = 1234, family = stats::gaussian,
    refresh = 0)
```

Back to broom

```
1  broom::tidy(bayes_tm)
```

Error in warn_on_stanreg(x): The supplied model object seems to be outputted from the rstanarm package. Tidiers for mixed model output now live in the broom.mixed package.

```
1  broom.mixed::tidy(bayes_tm)
```

```
# A tibble: 4 × 3
  term                  estimate std.error
  <chr>                    <dbl>     <dbl>
1 (Intercept)              95.4      63.9
2 volume                    0.828     0.0759
3 coverpaperback           -0.263     3.63
4 volume:coverpaperback    -0.197     0.0518
```

```
1  broom.mixed::glance(bayes_tm)
```

```
# A tibble: 1 × 4
  algorithm   pss  nobs sigma
  <chr>      <dbl> <int> <dbl>
1 sampling    4000    15  85.5
```

Augment

```
1  augment(bayes_tm, new_data=books)
```

```
# A tibble: 15 × 5
   .pred  .resid volume weight cover
   <dbl>   <dbl>  <dbl>  <dbl> <fct>
 1  829.  -28.6     885    800 hardback
 2  937.   12.9    1016    950 hardback
 3 1027.   22.6    1125   1050 hardback
 4  294.   56.3     239    350 hardback
 5  676.   73.7     701    750 hardback
 6  627.  -26.6     641    600 hardback
 7 1113.  -37.7    1228   1075 hardback
 8  353. -103.      412    250 paperback
 9  696.    4.34    953    700 paperback
10  680.  -30.5     929    650 paperback
11 1037.  -61.6    1492    975 paperback
```

Predictions

```r
bayes_tm |>
  augment(
    new_data = tidyr::expand_grid(
      volume = seq(0, 1500, by=5),
      cover = c("hardback", "paperback") |> as.factor()
    )
  ) |>
  rename(weight = .pred) |>
ggplot(aes(x = volume, y = weight, color = cover, group = cover)) +
  geom_line() +
  geom_point(data = books)
```

Performance

```r
1  lm_tm |>
2    augment(new_data = books) |>
3    yardstick::rmse(weight, .pred)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>          <dbl>
1 rmse    standard        68.9
```

```r
1  bayes_tm |>
2    augment(new_data = books) |>
3    yardstick::rmse(weight, .pred)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>          <dbl>
1 rmse    standard        72.0
```

Cross validation and Feature engineering

The Office & IMDB

```
1  (office_ratings = read_csv("data/office_ratings.csv"))
```

```
# A tibble: 188 × 6
   season episode title            imdb_rating total_votes air_date
    <dbl>   <dbl> <chr>                   <dbl>       <dbl> <date>
 1      1       1 Pilot                     7.6        3706 2005-03-24
 2      1       2 Diversity Day             8.3        3566 2005-03-29
 3      1       3 Health Care               7.9        2983 2005-04-05
 4      1       4 The Alliance              8.1        2886 2005-04-12
 5      1       5 Basketball                8.4        3179 2005-04-19
 6      1       6 Hot Girl                  7.8        2852 2005-04-26
 7      2       1 The Dundies               8.7        3213 2005-09-20
 8      2       2 Sexual Harassment         8.2        2736 2005-09-27
 9      2       3 Office Olympics           8.4        2742 2005-10-04
10      2       4 The Fire                  8.4        2713 2005-10-11
# i 178 more rows
```

Rating vs Air Date

Test-train split

```r
set.seed(123)
(office_split = initial_split(office_ratings, prop = 0.8))
```

```
<Training/Testing/Total>
<150/38/188>
```

```r
(office_train = training(office_split))
```

```
# A tibble: 150 × 6
   season episode title     imdb_rating total_votes
    <dbl>   <dbl> <chr>           <dbl>       <dbl>
 1      8      18 Last Da…          7.8        1429
 2      9      14 Vandali…          7.6        1402
 3      2       8 Perform…          8.2        2416
 4      9       5 Here Co…          7.1        1515
 5      3      22 Beach G…          9.1        2783
 6      7       1 Nepotism          8.4        1897
 7      3      15 Phyllis…          8.3        2283
 8      9      21 Livin' …          8.9        2041
 9      9      18 Promos            8          1445
10      8      12 Pool Pa…          8          1612
# i 140 more rows
# i 1 more variable: air_date <date>
```

```r
(office_test = testing(office_split))
```

```
# A tibble: 38 × 6
   season episode title    imdb_rating total_votes
    <dbl>   <dbl> <chr>          <dbl>       <dbl>
 1      1       2 Diversi…         8.3        3566
 2      2       4 The Fire         8.4        2713
 3      2       9 E-Mail …         8.4        2527
 4      2      12 The Inj…         9           3282
 5      2      22 Casino …         9.3        3644
 6      3       5 Initiat…         8.2        2254
 7      3      16 Busines…         8.8        2622
 8      3      17 Cocktai…         8.5        2264
 9      4       6 Branch …         8.5        2185
10      4       7 Survivo…         8.3        2110
# i 28 more rows
# i 1 more variable: air_date <date>
```

Feature engineering with dplyr

```r
1  office_train |>
2    mutate(
3      season = as_factor(season),
4      month = lubridate::month(air_date),
5      wday = lubridate::wday(air_date),
6      top10_votes = as.integer(total_votes > quantile(total_votes, 0.9))
7    )
```

```
# A tibble: 150 × 9
   season episode title              imdb_rating total_votes air_date    month  wday top10_votes
   <fct>    <dbl> <chr>                    <dbl>       <dbl> <date>      <dbl> <dbl>       <int>
 1 8           18 Last Day in Florida        7.8        1429 2012-03-08      3     5           0
 2 9           14 Vandalism                  7.6        1402 2013-01-31      1     5           0
 3 2            8 Performance Review         8.2        2416 2005-11-15     11     3           0
 4 9            5 Here Comes Treble          7.1        1515 2012-10-25     10     5           0
 5 3           22 Beach Games                9.1        2783 2007-05-10      5     5           0
 6 7            1 Nepotism                   8.4        1897 2010-09-23      9     5           0
 7 3           15 Phyllis' Wedding           8.3        2283 2007-02-08      2     5           0
 8 9           21 Livin' the Dream           8.9        2041 2013-05-02      5     5           0
 9 9           18 Promos                     8          1445 2013-04-04      4     5           0
10 8           12 Pool Party                 8          1612 2012-01-19      1     5           0
# i 140 more rows
```

Anyone see a potential problem with the code above?

recipes

www.rstudio.com

Better living through recipes

```r
1  r = recipe(imdb_rating ~ ., data = office_train)
```

```r
1  summary(r)
```

```
# A tibble: 6 × 4
  variable    type      role      source
  <chr>       <list>    <chr>     <chr>
1 season      <chr [2]> predictor original
2 episode     <chr [2]> predictor original
3 title       <chr [3]> predictor original
4 total_votes <chr [2]> predictor original
5 air_date    <chr [1]> predictor original
6 imdb_rating <chr [2]> outcome   original
```

Recipe roles

```r
1  r = recipe(
2    imdb_rating ~ ., data = office_train
3  ) |>
4    update_role(title, new_role = "ID")
```

```r
1  summary(r)
```

```
# A tibble: 6 × 4
  variable    type      role      source
  <chr>       <list>    <chr>     <chr>
1 season      <chr [2]> predictor original
2 episode     <chr [2]> predictor original
3 title       <chr [3]> ID        original
4 total_votes <chr [2]> predictor original
5 air_date    <chr [1]> predictor original
6 imdb_rating <chr [2]> outcome   original
```

Adding features (month & day of week)

```r
1  r = recipe(
2    imdb_rating ~ ., data = office_train
3  ) |>
4    update_role(title, new_role = "ID") |>
5    step_date(air_date, features = c("dow", "month"))
```

```r
1  summary(r)
```

```
# A tibble: 6 × 4
  variable    type      role      source
  <chr>       <list>    <chr>     <chr>
1 season      <chr [2]> predictor original
2 episode     <chr [2]> predictor original
3 title       <chr [3]> ID        original
4 total_votes <chr [2]> predictor original
5 air_date    <chr [1]> predictor original
6 imdb_rating <chr [2]> outcome   original
```

Adding Holidays

```r
r = recipe(
  imdb_rating ~ ., data = office_train
) |>
  update_role(title, new_role = "ID") |>
  step_date(air_date, features = c("dow", "month")) |>
  step_holiday(
    air_date,
    holidays = c("USThanksgivingDay", "USChristmasDay", "USNewYearsDay", "USIndependenceDay"),
    keep_original_cols = FALSE
  )
```

```r
summary(r)
```

```
# A tibble: 6 × 4
  variable     type       role      source
  <chr>        <list>     <chr>     <chr>
1 season       <chr [2]>  predictor original
2 episode      <chr [2]>  predictor original
3 title        <chr [3]>  ID        original
4 total_votes  <chr [2]>  predictor original
5 air_date     <chr [1]>  predictor original
6 imdb_rating  <chr [2]>  outcome   original
```

Seasons as factors

```r
r = recipe(
  imdb_rating ~ ., data = office_train
) |>
  update_role(title, new_role = "ID") |>
  step_date(air_date, features = c("dow", "month")) |>
  step_holiday(
    air_date,
    holidays = c("USThanksgivingDay", "USChristmasDay", "USNewYearsDay", "USIndependenceDay"),
    keep_original_cols = FALSE
  ) |>
  step_num2factor(season, levels = as.character(1:9))
```

```r
summary(r)
```

```
# A tibble: 6 × 4
  variable    type      role      source
  <chr>       <list>    <chr>     <chr>
1 season      <chr [2]> predictor original
2 episode     <chr [2]> predictor original
3 title       <chr [3]> ID        original
4 total_votes <chr [2]> predictor original
5 air_date    <chr [1]> predictor original
6 imdb_rating <chr [2]> outcome   original
```

Dummy coding

```r
r = recipe(
  imdb_rating ~ ., data = office_train
) |>
  update_role(title, new_role = "ID") |>
  step_date(air_date, features = c("dow", "month")) |>
  step_holiday(
    air_date,
    holidays = c("USThanksgivingDay", "USChristmasDay", "USNewYearsDay", "USIndependenceDay"),
    keep_original_cols = FALSE
  ) |>
  step_num2factor(season, levels = as.character(1:9)) |>
  step_dummy(all_nominal_predictors())
```

```r
summary(r)
```

```
# A tibble: 6 × 4
  variable     type       role      source
  <chr>        <list>     <chr>     <chr>
1 season       <chr [2]>  predictor original
2 episode      <chr [2]>  predictor original
3 title        <chr [3]>  ID        original
4 total_votes  <chr [2]>  predictor original
5 air_date     <chr [1]>  predictor original
6 imdb_rating  <chr [2]>  outcome   original
```

top10_votes

```r
r = recipe(
  imdb_rating ~ ., data = office_train
) |>
  update_role(title, new_role = "ID") |>
  step_date(air_date, features = c("dow", "month")) |>
  step_holiday(
    air_date,
    holidays = c("USThanksgivingDay", "USChristmasDay", "USNewYearsDay", "USIndependenceDay"),
    keep_original_cols = FALSE
  ) |>
  step_num2factor(season, levels = as.character(1:9)) |>
  step_dummy(all_nominal_predictors()) |>
  step_percentile(total_votes) |>
  step_mutate(top10 = as.integer(total_votes >= 0.9)) |>
  step_rm(total_votes)
```

```r
summary(r)
```

```
# A tibble: 6 × 4
  variable    type      role      source
  <chr>       <list>    <chr>     <chr>
1 season      <chr [2]> predictor original
2 episode     <chr [2]> predictor original
3 title       <chr [3]> ID        original
4 total_votes <chr [2]> predictor original
```

Preparing a recipe

```
1  prep(r)
```

```
── Recipe ──────────────────────────────────────────

── Inputs

Number of variables by role

outcome:    1
predictor:  4
ID:         1

── Training information

Training data contained 150 data points and no incomplete rows.

── Operations

• Date features from: air_date | Trained

• Holiday features from: air_date | Trained

• Factor variables from: season | Trained

• Dummy variables from: season, air_date_dow, air_date_month | Trained

• Percentile transformation on: total_votes | Trained

• Variable mutation for: ~as.integer(total_votes >= 0.9) | Trained

• Variables removed: total_votes | Trained
```

Baking a recipe

```
1  prep(r) |>
2    bake(new_data = office_train)
```

```
# A tibble: 150 × 33
   episode title     imdb_rating air_date_USThanksgiv…¹ air_date_USChristmas…² air_date_USNewYearsDay
     <dbl> <fct>           <dbl>                  <int>                  <int>                  <int>
 1      18 Last Da…          7.8                      0                      0                      0
 2      14 Vandali…          7.6                      0                      0                      0
 3       8 Perform…          8.2                      0                      0                      0
 4       5 Here Co…          7.1                      0                      0                      0
 5      22 Beach G…          9.1                      0                      0                      0
 6       1 Nepotism          8.4                      0                      0                      0
 7      15 Phyllis…          8.3                      0                      0                      0
 8      21 Livin' …          8.9                      0                      0                      0
 9      18 Promos            8                        0                      0                      0
10      12 Pool Pa…          8                        0                      0                      0
# ℹ 140 more rows
# ℹ abbreviated names: ¹air_date_USThanksgivingDay, ²air_date_USChristmasDay
# ℹ 27 more variables: air_date_USIndependenceDay <int>, season_X2 <dbl>, season_X3 <dbl>,
#   season_X4 <dbl>, season_X5 <dbl>, season_X6 <dbl>, season_X7 <dbl>, season_X8 <dbl>,
```

Informative features?

```
1  prep(r) |>
2    bake(new_data = office_train) |>
3    map_int(~ length(unique(.x)))
```

episode	title	imdb_rating
26	150	26
air_date_USThanksgivingDay	air_date_USChristmasDay	air_date_USNewYearsDay
1	1	1
air_date_USIndependenceDay	season_X2	season_X3
1	2	2
season_X4	season_X5	season_X6
2	2	2
season_X7	season_X8	season_X9
2	2	2
air_date_dow_Mon	air_date_dow_Tue	air_date_dow_Wed
1	2	1
air_date_dow_Thu	air_date_dow_Fri	air_date_dow_Sat
2	1	1
air_date_month_Feb	air_date_month_Mar	air_date_month_Apr
2	2	2
air_date_month_May	air_date_month_Jun	air_date_month_Jul

Removing zero variance predictors

```r
1  r = recipe(
2      imdb_rating ~ ., data = office_train
3  ) |>
4  update_role(title, new_role = "ID") |>
5  step_date(air_date, features = c("dow", "month")) |>
6  step_holiday(
7    air_date,
8    holidays = c("USThanksgivingDay", "USChristmasDay", "USNewYearsDay", "USIndependenceDay"),
9    keep_original_cols = FALSE
10 ) |>
11 step_num2factor(season, levels = as.character(1:9)) |>
12 step_dummy(all_nominal_predictors()) |>
13 step_percentile(total_votes) |>
14 step_mutate(top10 = as.integer(total_votes >= 0.9)) |>
15 step_rm(total_votes) |>
16 step_zv(all_predictors())
```

```
1  prep(r) |>
2    bake(new_data = office_train)
```

A tibble: 150 × 22
   episode title   imdb_rating season_X2 season_X3 season_X4 season_X5 season_X6 season_X7 season_X8
     <dbl> <fct>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
 1   18 Last D…   7.8   0   0   0   0   0   0   1
 2   14 Vandal…   7.6   0   0   0   0   0   0   0
 3   8 Perfor…   8.2   1   0   0   0   0   0   0
 4   5 Here C…   7.1   0   0   0   0   0   0   0
 5   22 Beach …   9.1   0   1   0   0   0   0   0
 6   1 Nepoti…   8.4   0   0   0   0   0   1   0
 7   15 Phylli…   8.3   0   1   0   0   0   0   0
 8   21 Livin'…   8.9   0   0   0   0   0   0   0
 9   18 Promos   8   0   0   0   0   0   0   0
10   12 Pool P…   8   0   0   0   0   0   0   1
ℹ 140 more rows
ℹ 12 more variables: season_X9 <dbl>, air_date_dow_Tue <dbl>, air_date_dow_Thu <dbl>,
air_date_month_Feb <dbl>, air_date_month_Mar <dbl>, air_date_month_Apr <dbl>,
air_date_month_May <dbl>, air_date_month_Sep <dbl>, air_date_month_Oct <dbl>,

Really putting it all together

```
1  (office_work = workflow() |>
2    add_recipe(r) |>
3    add_model(
4      linear_reg() |>
5      set_engine("lm")
6    )
7  )
```

```
══ Workflow ════════════════════════════════════
Preprocessor: Recipe
Model: linear_reg()


── Preprocessor ─────────────────────────────────

8 Recipe Steps

• step_date()
• step_holiday()
• step_num2factor()
• step_dummy()
• step_percentile()
• step_mutate()
• step_rm()
• step_zv()

── Model ─────────────────────────────────────────
```

Workflow fit

```
1  (office_fit = office_work |>
2     fit(data = office_train))
```

```
══ Workflow [trained] ══════════════════════════════════════════
Preprocessor: Recipe
Model: linear_reg()

── Preprocessor ────────────────────────────────────────────────
8 Recipe Steps

• step_date()
• step_holiday()
• step_num2factor()
• step_dummy()
• step_percentile()
• step_mutate()
• step_rm()
• step_zv()

── Model ───────────────────────────────────────────────────────
```

Performance

```
1  office_fit |>
2    augment(office_train) |>
3    rmse(imdb_rating, .pred)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>          <dbl>
1 rmse    standard       0.342
```

```
1  office_fit |>
2    augment(office_test) |>
3    rmse(imdb_rating, .pred)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>          <dbl>
1 rmse    standard       0.399
```

k-fold cross validation

Creating folds

```r
1  set.seed(123)
2  (folds = vfold_cv(office_train, v=5))
```

```
#  5-fold cross-validation
# A tibble: 5 × 2
  splits            id
  <list>            <chr>
1 <split [120/30]>  Fold1
2 <split [120/30]>  Fold2
3 <split [120/30]>  Fold3
4 <split [120/30]>  Fold4
5 <split [120/30]>  Fold5
```

```r
1  (office_fit_folds = office_work |>
2    fit_resamples(folds)
3  )
```

```
→ A | warning: prediction from rank-deficient fit;
consider predict(., rankdeficient="NA")

There were issues with some computations   A: x1
There were issues with some computations   A: x1

# Resampling results
# 5-fold cross-validation
# A tibble: 5 × 4
  splits            id     .metrics
  <list>            <chr> <list>
1 <split [120/30]>  Fold1 <tibble [2 × 4]>
2 <split [120/30]>  Fold2 <tibble [2 × 4]>
3 <split [120/30]>  Fold3 <tibble [2 × 4]>
4 <split [120/30]>  Fold4 <tibble [2 × 4]>
5 <split [120/30]>  Fold5 <tibble [2 × 4]>
  .notes
  <list>
1 <tibble [0 × 3]>
2 <tibble [1 × 3]>
3 <tibble [0 × 3]>
4 <tibble [0 × 3]>
5 <tibble [0 × 3]>
```

Fold performance

```
1  tune::collect_metrics(office_fit_folds)
```

```
# A tibble: 2 × 6
  .metric .estimator  mean     n std_err .config
  <chr>   <chr>      <dbl> <int>   <dbl> <chr>
1 rmse    standard   0.420     5  0.0182 Preprocessor1_Model1
2 rsq     standard   0.429     5  0.0597 Preprocessor1_Model1
```

```
1  tune::collect_metrics(office_fit_folds, summarize = FALSE) |>
2    filter(.metric == "rmse")
```

```
# A tibble: 5 × 5
  id    .metric .estimator .estimate .config
  <chr> <chr>   <chr>          <dbl> <chr>
1 Fold1 rmse    standard       0.467 Preprocessor1_Model1
2 Fold2 rmse    standard       0.403 Preprocessor1_Model1
3 Fold3 rmse    standard       0.405 Preprocessor1_Model1
4 Fold4 rmse    standard       0.454 Preprocessor1_Model1
5 Fold5 rmse    standard       0.368 Preprocessor1_Model1
```

More on the tune package next time